

Machine Learning Canvas

	PREDICTIONS	OBJECTIVES	DATA
BACKGROUND	End-user Who will use the predictive application/who will be affected by it? <div></div>	Value proposition What are we trying to do for the system's users? (e.g. spend less time on X, increase Y...) <div>Hahhaha</div>	Data sources Where do/can we get data from? (internal database, 3rd party API, etc.) <div></div>
ENGINE SPECS	Problem Decision to be made (on behalf of user) <div></div> Input (i.e. question "parameter") <div></div> Possible outputs (i.e. "answers") <div></div> Applicable trade-offs (i.e. "time vs. quality") <div></div> Type of problem (e.g. classification, regression, recommendation...) <div></div> Baseline: simple, alternative way of making predictions (e.g. manual rules) <div></div>	Performance evaluation Domain-specific/bottom-line metrics for monitoring performance in production <div></div> Prediction accuracy metrics (e.g. MSE if regression; % accuracy, #FP for classification) <div></div> Offline performance evaluation method (e.g. cross-validation or simple training/test split) <div></div>	Data preparation How do we get training data (inputs, and outputs if supervised learning)? How many data points? <div></div> Input features (extracted from data sources). If too many, list types of features and mention key ones. <div></div>
INTEGRATION	Using predictions When do we make predictions and how many? <div></div> What is the time constraint for making those predictions? <div></div> How do we use predictions and confidence values? <div></div>		Learning models When do we create/update models? With what data/how much? <div></div> What is the time constraint for creating a model? <div></div> Criteria for deploying model (e.g. minimum performance value — absolute, relative to baseline or to previous model) <div></div>

Reset Form

Machine Learning Canvas v0.1

[Louis Dorard](#) © 2015. Please reference [machinelearningcanvas.com](#) by linking to it if you use the canvas.